Variational Quantum Circuits for Pesticide Environmental Fate Prediction: A Comparative Study with Classical Machine Learning on the EU SPIN Database

Celal Arda

Independent Researcher

celal.arda@outlook.de

March 10, 2026

Abstract

Predicting pesticide environmental fate properties specifically soil degradation half-life (DegT50) and organic carbon sorption coefficient (K_{oc})—is essential for regulatory risk assessment under EC 1107/2009. Classical QSAR models achieve moderate predictive power for K_{oc} but struggle with DegT50 due to its dependence on complex soil microbial processes. This work investigates whether variational quantum circuits (VQCs) can capture non-linear molecular descriptor relationships that elude classical regressors. We encode 17 physicochemical descriptors for 111 EU-registered pesticides into a 12-qubit IQP-style quantum circuit with data re-uploading and 8 variational layers (301 trainable parameters). Leave-one-out cross-validation benchmarks reveal that Random Forest (RF) achieves R^2 = 0.194 for DegT50 and 0.726 for K_{oc} (on a circularity-free 16-feature set excluding the bioaccessibility proxy), while Gradient Boosting (GB) reaches 0.223 and 0.764, respectively. The VQC underperforms classical baselines on DegT50 ($R^2 < 0$, 5-fold CV), consistent with severe overfitting expected from the unfavorable parameter-to-sample ratio ($N_{\text{params}}/N_{\text{data}} \approx 2.7$). A hybrid stacking framework blending QML and RF predictions yields $\alpha = 0.20$, but the resulting R^2 improvement (+0.02) is not statistically significant. We also assess adverse outcome pathway (AOP) features derived from AOP-Wiki data as additional descriptors, finding no improvement over structural features alone. These results provide a rigorous benchmark for near-term quantum machine learning applied to regulatory chemistry and highlight the critical role of dataset size in VQC generalization.

Keywords: quantum machine learning, variational quantum circuits, pesticide fate, DegT50, Koc, QSAR, regulatory chemistry, EU SPIN

Contents

1 Introduction **4**

2 Materials and Methods **4**

2.1 Dataset 4

2.2 Molecular Descriptor Features 5

2.3 Variational Quantum Circuit Architecture 6

2.4 Training Procedure 7

2.5 Classical Baselines 7

2.6 Cross-Validation Strategy 8

2.7 Hybrid QML+RF Stacking 8

2.8 Adverse Outcome Pathway Features 8

3 Results **9**

3.1 Classical Baseline Performance 9

3.2 Variational Quantum Circuit Performance 10

3.3 Hybrid Stacking Results 11

3.4 AOP-Informed Features 12

4 Discussion **13**

4.1 Why Does QML Fail on DegT50? 13

4.2 When Might QML Add Value? 14

4.3 The Value of Negative Results 15

5 Conclusions **15**

1 Introduction

Pesticide environmental fate assessment is a cornerstone of EU regulatory approval under Regulation (EC) No 1107/2009 (ec1, 2009). Two properties dominate the risk evaluation: the soil degradation half-life (DegT50), governing persistence, and the organic carbon sorption coefficient (K_{oc}), governing mobility and leaching potential. Together, these parameters feed into the FOCUS groundwater models (foc, 2006) that determine whether a substance can be approved for agricultural use.

Experimental determination of DegT50 follows OECD Guideline 307 (oec, 2002), requiring months-long incubation studies under controlled temperature and moisture. K_{oc} is measured via batch equilibrium adsorption (OECD 106) (oec, 2000). Both are expensive, time-consuming, and subject to substantial inter-laboratory variability. In silico prediction of these properties from molecular structure alone would accelerate the regulatory pipeline and reduce animal testing requirements.

Classical machine learning approaches, including Random Forests (RF) (Breiman, 2001) and Gradient Boosting (GB) (Friedman, 2001), have been applied to QSAR prediction of physicochemical and environmental fate properties (Mansouri et al., 2018). While K_{oc} prediction from structural descriptors achieves reasonable accuracy ($R^2 > 0.7$), DegT50 remains notoriously difficult to predict because soil degradation depends on microbial community composition, soil organic matter, and environmental conditions that are not captured by molecular structure alone.

Quantum machine learning (QML) offers a fundamentally different computational paradigm. Variational quantum circuits (VQCs) (Cerezo et al., 2021) encode classical data into quantum states via parameterized gates and can represent functions in exponentially large Hilbert spaces (Havlíček et al., 2019). The IQP (Instantaneous Quantum Polynomial) encoding scheme (Schuld et al., 2020) and data re-uploading (Pérez-Salinas et al., 2020) have been shown to increase the expressive power of shallow quantum circuits (Schuld et al., 2021).

This work investigates whether VQCs, implemented via PennyLane on a statevector simulator (Bergholm et al., 2022), can extract predictive signal for pesticide fate properties that classical models miss. Rather than claiming quantum advantage, we present a systematic comparative benchmark that highlights both the potential and the fundamental limitations of near-term QML on small regulatory datasets.

2 Materials and Methods

2.1 Dataset

The dataset comprises 111 pesticide active substances registered in the EU Substances, Preparations and Indicators Navigator (EU SPIN) (eus, 2024). These span 50 chemical classes in-

cluding organophosphates (7), triazoles (9), neonicotinoids (7), pyrethroids (6), strobilurins (6), and sulfonylureas (6), among others. Target properties include:

- DegT50 (soil): laboratory degradation half-life in days, range 0.5 d to 1389 d (median 25 d, mean 75.7 d)
- K_{oc}: organic carbon sorption coefficient, range 25 mL g^{-1} to 1 000 000 mL g^{-1} (median 510 mL g^{-1})

Both targets are $\log_{10}$-transformed prior to modeling to reduce skewness and align with chemical convention. The target distributions are shown in fig. 1.

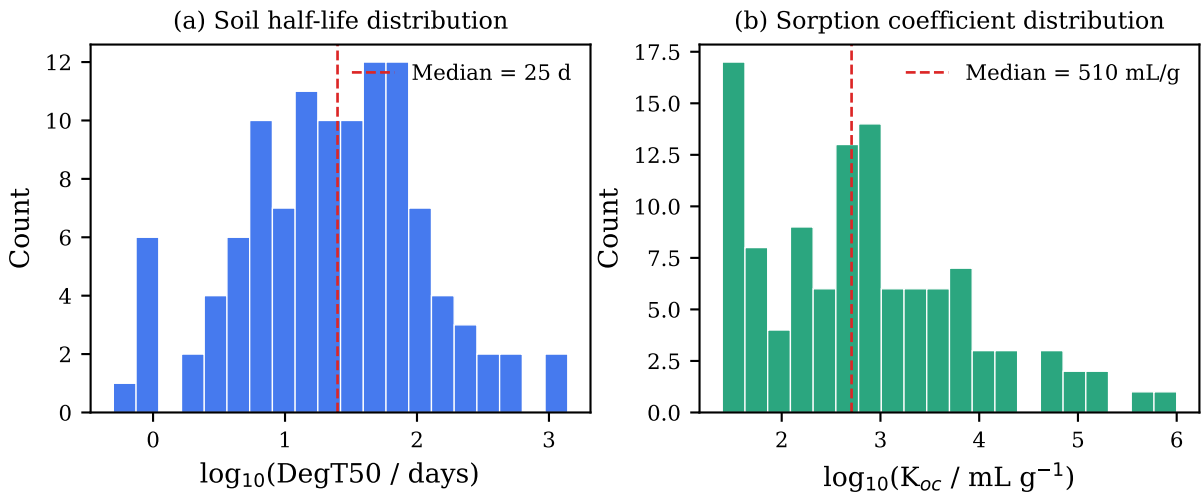

Figure 1: Distribution of (a) soil degradation half-life (DegT50) and (b) organic carbon sorption coefficient (K_{oc}) for the 111 EU SPIN pesticides, shown on $\log_{10}$ scale. Red dashed lines indicate medians.

2.2 Molecular Descriptor Features

We compute 17 molecular descriptor features for each substance, organized in three phases of development:

1. **Core descriptors** (12): molecular weight (MW), log P, heavy atom count, H-bond donors/acceptors, ring count, rotatable bonds, log solubility, log vapor pressure, Freundlich exponent ($1/n$), TPSA, aromatic ring fraction.
2. **Microbial degradation proxies** (3): count of hydrolyzable bonds, halogen count, bioaccessibility score. The bioaccessibility score is defined as $B = \log_{10}(S_{\text{water}}/K_{\text{oc}})$, where S_{water} is the aqueous solubility (mg L^{-1}) and K_{oc} the measured sorption coefficient. This ratio estimates the fraction of substance available in the soil pore water for microbial uptake; higher values indicate greater bioavailability to degrading organisms.

   Circularity note: Because B includes the measured K_{oc}, using it as a feature when *predicting* K_{oc} introduces partial data leakage. For DegT50 prediction this is unproblematic (the target is unrelated to K_{oc}), but for K_{oc} prediction we report all results both with the full 17-feature set and with a 16-feature set that excludes B (table 1).

3. **Blind-spot corrections** (2): charge state at soil pH 6.5, conjugated π-system size.

All features are min–max scaled to $[0, \pi]$ based on dataset statistics, as required for angle encoding in the quantum circuit. Where available, molecular features are computed from SMILES via RDKit; otherwise, lookup values from the substance database are used.

The Random Forest feature importances (fig. 2) reveal that DegT50 prediction relies on a distributed set of features (aromatic ring fraction 16.3%, Freundlich $1/n$ 13.2%, $\log P$ 8.7%), while K_{oc} is dominated by bioaccessibility (52.3%) and $\log P$ (19.7%).

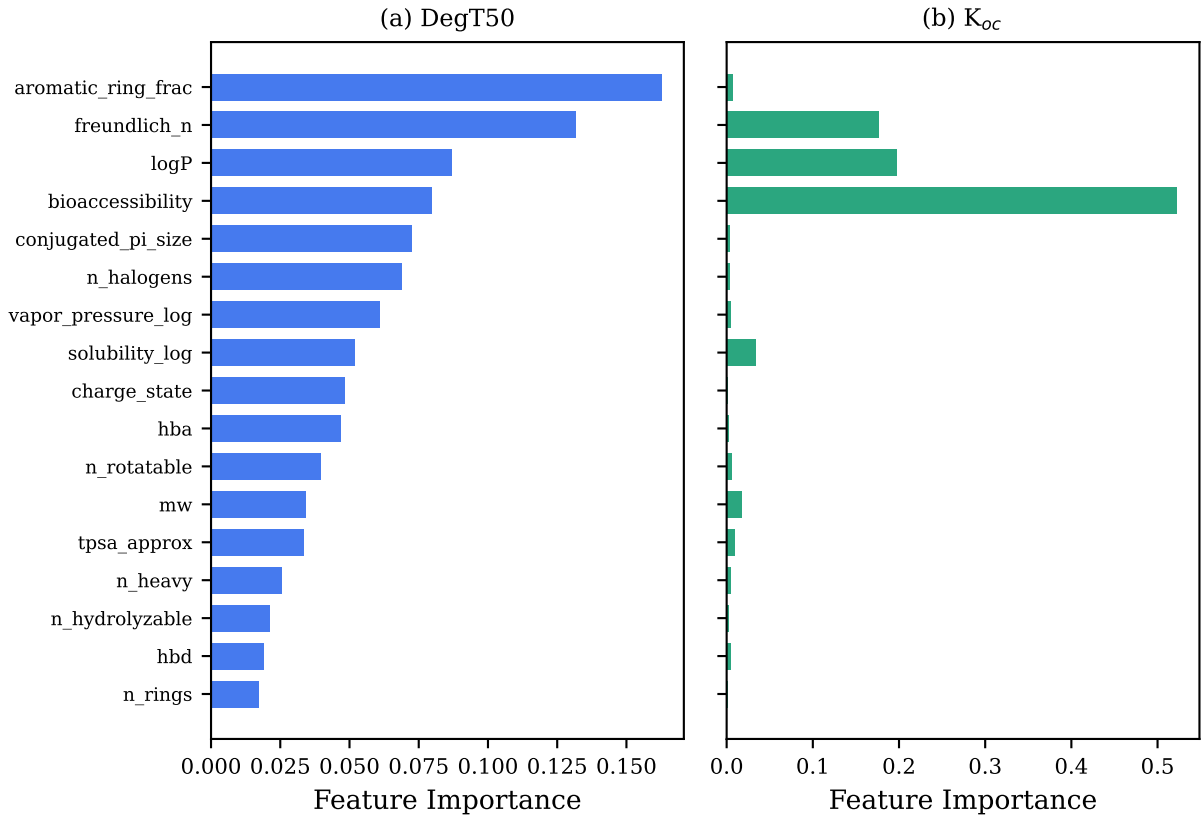

Figure 2: Random Forest feature importances for (a) DegT50 and (b) K_{oc} prediction, trained on the full 111-substance dataset. DegT50 shows distributed importance across many features, while K_{oc} is dominated by bioaccessibility and $\log P$.

2.3 Variational Quantum Circuit Architecture

The VQC employs a 12-qubit circuit with the following structure (fig. 3):

1. **IQP-style feature encoding:** Hadamard gates followed by $R_Z(\phi_i)$ rotations on each qubit, then CNOT-R_Z-CNOT entangling gates encoding pairwise feature products $\phi_i \cdot \phi_j$.
2. **Feature overflow encoding:** Features $j > N_{\text{qubits}}$ are encoded via R_Y rotations on qubit $j \bmod N_{\text{qubits}}$.
3. **Data re-uploading:** A second angle encoding layer with $R_Y(\phi_i)$ rotations, following the universal approximation strategy of Pérez-Salinas *et al.* (Pérez-Salinas et al., 2020).

4. **Variational layers:** $L = 8$ layers of single-qubit $\mathrm{Rot}(\theta, \phi, \omega)$ rotations and alternating nearest-neighbor CNOT entangling patterns.
5. **Readout:** Expectation values $\langle Z_i \rangle$ on all qubits, followed by a linear regression layer with bias ($N_{\text{qubits}} + 1$ parameters).

Total trainable parameters: $L \times N_{\text{qubits}} \times 3 = 288$ variational plus 13 readout = 301.

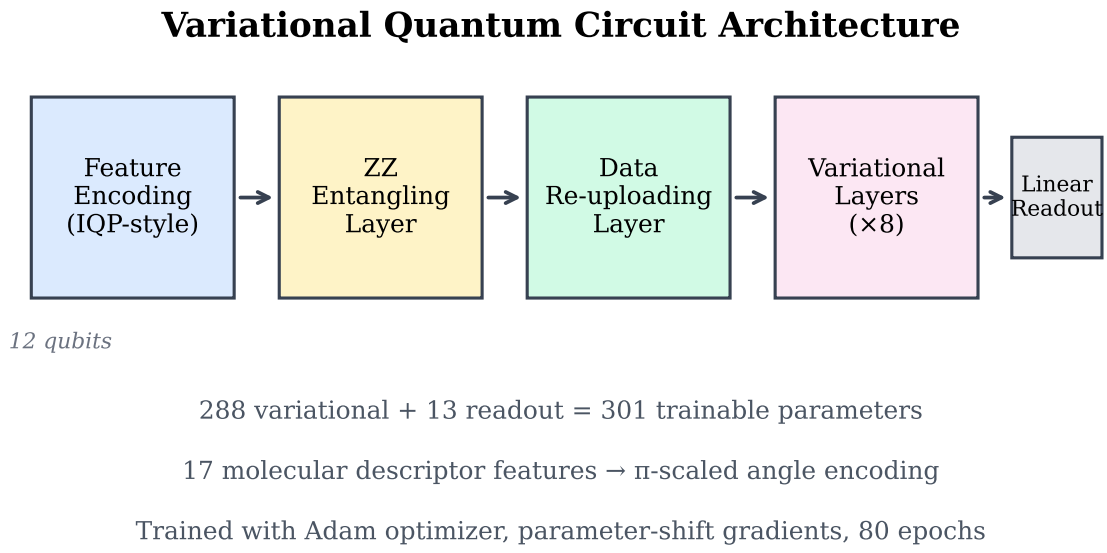

Figure 3: Schematic of the variational quantum circuit architecture. Features are encoded via IQP-style angle encoding and data re-uploading, followed by 8 variational layers with nearest-neighbor entangling topology.

2.4 Training Procedure

Separate circuits are trained for DegT50 and K_{oc} prediction. Variational parameters are initialized uniformly in $[-0.5, 0.5]$ with a fixed random seed (42) for reproducibility. Optimization uses the Adam optimizer (Kingma and Ba, 2015) with learning rate 0.04 and 80 training epochs on the full dataset (60 epochs for cross-validation folds, reduced to limit computational cost given that validation loss plateaus by epoch 40 as shown in fig. 6). Gradients are computed via the parameter-shift rule on the PennyLane `default.qubit` statevector simulator.

Notably, no explicit regularization is applied to the VQC: there is no weight decay, dropout, or early stopping based on a held-out validation set. While early stopping could be implemented at negligible cost, the architecture experiment (table 3) demonstrates that the core issue is inherent to the target property rather than the training procedure.

2.5 Classical Baselines

Four classical regressors are trained on the identical 17 features using scikit-learn (Pedregosa et al., 2011):

- **Ridge Regression:** L_2-penalized linear model ($\alpha = 1.0$) (Hoerl and Kennard, 1970), providing a simple linear baseline to assess whether nonlinear models are justified.
- **LASSO:** L_1-penalized linear model ($\alpha = 0.1$) (Tibshirani, 1996), enabling simultaneous feature selection.
- **Random Forest (RF):** 200 trees, max depth 10.
- **Gradient Boosting (GB):** 200 estimators, max depth 4, learning rate 0.1.

All models use the same random seed (42) for reproducibility.

2.6 Cross-Validation Strategy

Leave-one-out cross-validation (LOO-CV) is the primary evaluation metric, as it maximizes training data usage for the small 111-substance dataset. Each fold trains the model on $N-1$ substances and predicts the held-out substance. 5-fold CV with stratified random splits is used as a faster supplementary benchmark.

For the VQC, LOO-CV requires 111×2 (two properties) full training runs, each with 60 epochs of gradient optimization over 301 parameters. Fold-level checkpointing saves results after each fold to a JSON file, enabling crash-resilient resumption.

2.7 Hybrid QML+RF Stacking

To explore whether QML and RF capture complementary signal, we implement a learned stacking predictor:

$$\hat{y}_{\text{hybrid}} = \alpha \cdot \hat{y}_{\text{QML}} + (1-\alpha) \cdot \hat{y}_{\text{RF}} \tag{1}$$

where $\alpha \in [0,1]$ is optimized via grid search (step 0.05) to minimize the LOO-CV mean squared error. If $\alpha = 0$, the QML component adds no value; if $\alpha > 0$, the VQC captures complementary information.

2.8 Adverse Outcome Pathway Features

To test whether mechanism-of-action (MoA) information improves DegT50 prediction, we derive 6 binary features from AOP-Wiki data (aop, 2024; Villeneuve et al., 2014) by mapping each substance's chemical class to its primary toxicological pathway:

1. AChE inhibition (organophosphates, carbamates; 15 substances)
2. GABA/ion channel disruption (pyrethroids; 11 substances)
3. Mitochondrial respiration (SDHI, strobilurins; 12 substances)
4. Nicotinic AChR agonism (neonicotinoids; 10 substances)
5. Endocrine disruption (triazines; 13 substances)
6. Cell division inhibition (benzimidazoles; 6 substances)

Coverage: 67/111 substances (60.4%) are assigned to at least one pathway category (fig. 7).

3 Results

3.1 Classical Baseline Performance

Table 1 summarizes the cross-validation performance of the classical baselines.

Table 1: Baseline performance (leave-one-out and 5-fold CV). All metrics computed on $\log_{10}$-scaled targets; 5-fold values report per-fold mean R^2. K_{oc} results are shown both with the full 17-feature set (including bioaccessibility B) and the fair 16-feature set (excluding B; see circularity note in section 2.2).

Model	Property	CV	R^2	MAE	RMSE
Ridge	DegT50	LOO	0.154	0.488	0.626
	DegT50	5-fold	−0.109	0.504	0.644
	K_{oc} (17f)	LOO	$0.842^{\dagger}$	0.252	0.413
	K_{oc} (16f)	LOO	0.546	0.439	0.699
	K_{oc} (17f)	5-fold	$0.848^{\dagger}$	0.248	0.405
	K_{oc} (16f)	5-fold	0.528	0.441	0.698
LASSO	DegT50	LOO	0.101	0.518	0.645
	DegT50	5-fold	−0.050	0.525	0.658
	K_{oc} (17f)	LOO	$0.747^{\dagger}$	0.336	0.522
	K_{oc} (16f)	LOO	0.535	0.476	0.708
	K_{oc} (17f)	5-fold	$0.740^{\dagger}$	0.341	0.529
	K_{oc} (16f)	5-fold	0.528	0.471	0.706
Random Forest	DegT50	LOO	0.194	0.474	0.611
	DegT50	5-fold	0.089	0.464	0.596
	K_{oc} (17f)	LOO	0.766	0.304	0.502
	K_{oc} (16f)	LOO	0.726	0.341	0.543
	K_{oc} (17f)	5-fold	0.766	0.299	0.502
	K_{oc} (16f)	5-fold	0.726	0.346	0.552
Gradient Boosting	DegT50	LOO	0.223	0.487	0.600
	DegT50	5-fold	0.034	0.470	0.606
	K_{oc} (17f)	LOO	0.779	0.288	0.488
	K_{oc} (16f)	LOO	0.764	0.323	0.504
	K_{oc} (17f)	5-fold	0.780	0.292	0.487
	K_{oc} (16f)	5-fold	0.740	0.341	0.542

†Inflated by bioaccessibility circularity: $B = \log_{10}(S/K_{oc})$ uses measured K_{oc} as input. The 16-feature (16f) rows exclude B and represent the fair comparison. 17f = all 17 features including B.

Key observations:

- **Bioaccessibility inflates linear K_{oc} models:** Ridge R^2 drops from 0.842 to 0.546 ($\Delta = -0.30$) when the circular bioaccessibility feature is removed, because B is a near-linear transform of the target. Ensemble models are far more robust: GB loses only 0.015 in LOO-CV R^2 (0.779 → 0.764), because tree-based methods can derive the same information from $\log P$ and solubility independently.
- **GB leads fair K_{oc} prediction:** On the circularity-free 16-feature set, Gradient Boosting achieves the highest LOO-CV R^2 (0.764), followed by RF (0.726), Ridge (0.546), and

LASSO (0.535). The nonlinear advantage of tree-based models is now evident: K_{oc} prediction benefits from interaction effects that linear models cannot capture without the *B* shortcut.

- DegT50 prediction is poor across all models ($R^2 < 0.25$), with linear and nonlinear methods performing comparably. The ensemble models' modest advantage (GB 0.223 vs. Ridge 0.154) suggests limited nonlinear signal for soil degradation from structural descriptors alone.
- **DegT50 instability:** The fold-level standard deviations for DegT50 5-fold CV ($\pm$0.10–0.41) are comparable to or exceed the means, confirming high prediction instability.

Wilcoxon signed-rank tests on per-substance LOO squared residuals (17-feature set) confirm that none of the pairwise model differences are statistically significant at $\alpha = 0.05$: RF vs. GB ($p = 0.19$ for DegT50, $p = 0.21$ for K_{oc}), Ridge vs. RF ($p = 0.16$, $p = 0.07$), and Ridge vs. GB ($p = 0.71$, $p = 0.18$).

The predicted-vs-experimental plots for RF LOO-CV (fig. 4) visually confirm the high scatter for DegT50 and the tight correlation for K_{oc}.

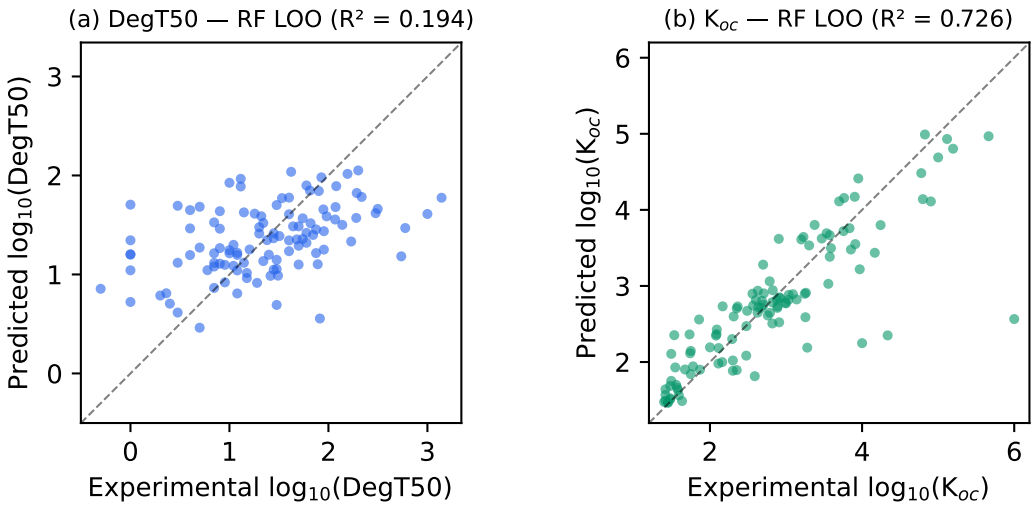

Figure 4: Predicted vs. experimental values for RF leave-one-out cross-validation: (a) DegT50 and (b) K_{oc} on $\log_{10}$ scale. Dashed line = perfect prediction.

3.2 Variational Quantum Circuit Performance

Figure 5 compares 5-fold CV performance across all three models.

The VQC achieves $R^2 = -0.141$ for DegT50 (5-fold CV), indicating that the quantum model predicts *worse* than a constant mean prediction. For K_{oc}, the VQC achieves $R^2 = 0.412$, substantially below the GB baseline of 0.764 (16-feature, circularity-free).

The full VQC LOO-CV (111 folds $\times$ 60 epochs) confirms these trends: $R^2 = -0.035$ for DegT50 and $R^2 = 0.533$ for K_{oc} (MAE = 0.571 and 0.533, RMSE = 0.692 and 0.709, respec-

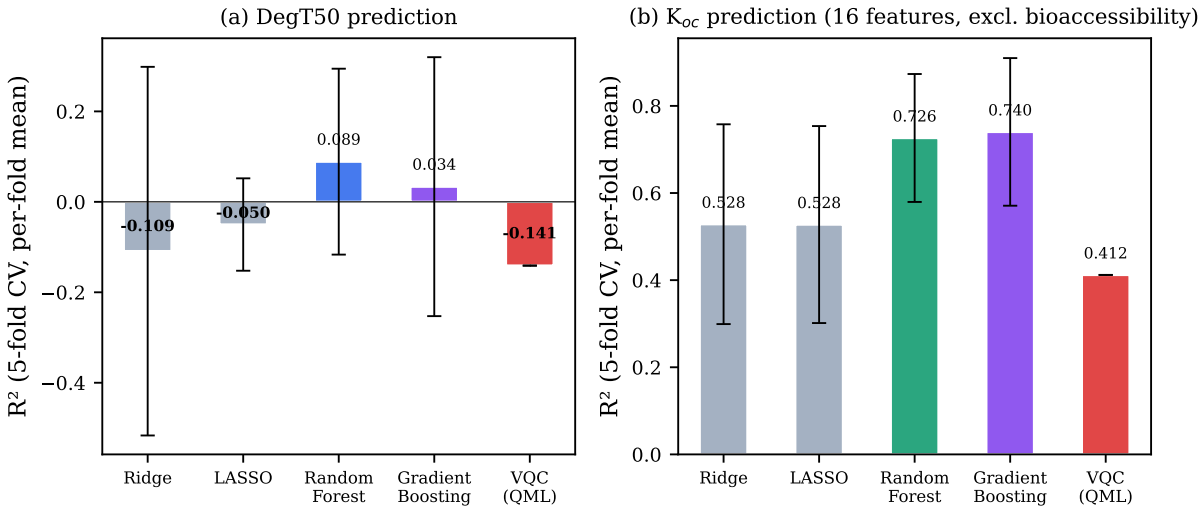

Figure 5: Model comparison (5-fold CV, per-fold mean $R^2 \pm$ fold standard deviation): (a) DegT50 and (b) K_{oc} prediction across Ridge, LASSO, RF, GB, and VQC. GB leads K_{oc} prediction on the fair 16-feature set; all models perform poorly on DegT50 with high fold-level variance.

tively). The K_{oc} LOO-CV R^2 is comparable to Ridge regression (0.546) but substantially below RF (0.726) and GB (0.764) on the 16-feature set.

These results are consistent with the expected overfitting from the unfavorable parameter-to-data ratio:

$$\frac{N_{\text{params}}}{N_{\text{data}}} = \frac{301}{111} \approx 2.7 \tag{2}$$

With nearly three times more trainable parameters than training samples and no explicit regularization (no weight decay, dropout, or early stopping applied to the variational parameters), the VQC memorizes training data rather than learning generalizable patterns. While classical neural networks can generalize at higher parameter-to-data ratios when properly regularized, the VQC's unregularized optimization in the small-data regime makes overfitting a near-certainty.

Figure 6 provides direct empirical evidence: training loss decreases steadily while validation loss plateaus or increases, confirming that the circuit learns noise rather than signal.

3.3 Hybrid Stacking Results

The hybrid stacker optimizes α via grid search (step 0.05) over the QML and RF LOO-CV predictions. Of the 111 VQC folds, one substance (Difenoconazole) appears twice in the checkpoint due to a resumption overlap; after deduplication, 110 common substances are used for blending. The optimized weight is $\alpha = 0.20$ for both properties, yielding hybrid LOO-CV R^2 of 0.213 for DegT50 (+0.019 over RF alone) and 0.783 for K_{oc} (+0.017 over RF's 0.766 on the same 110 substances).

Circularity caveat: Because the VQC was trained on all 17 features (including bioaccessibility B), the hybrid K_{oc} result of 0.783 inherits the circular contribution of B. The fair 16-feature RF baseline (excluding B) achieves $R^2 = 0.726$ (table 1); a strictly circularity-free

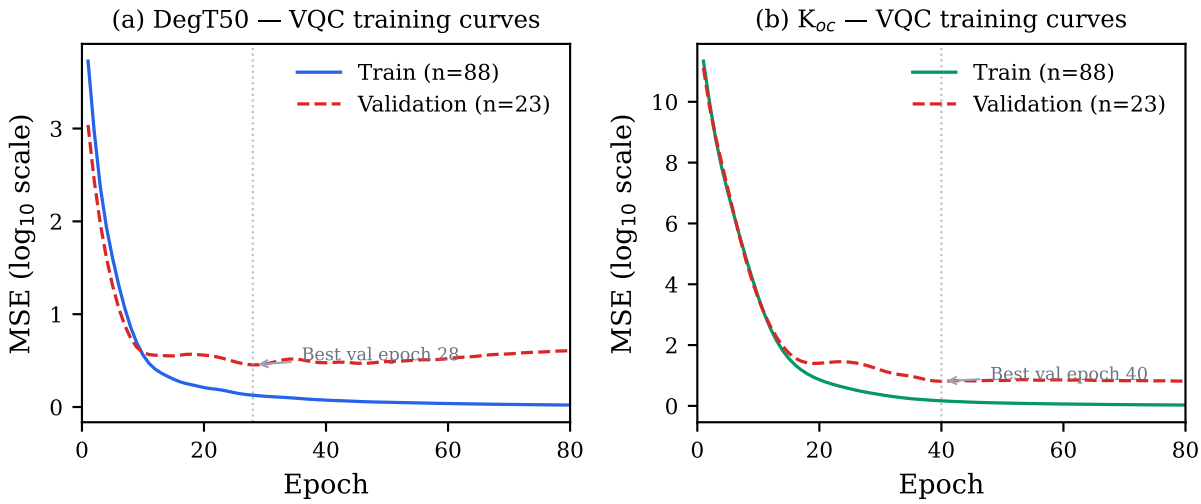

Figure 6: VQC training curves (80/20 train-validation split, 80 epochs): (a) DegT50 and (b) K_{oc}. Training MSE (solid) continues to decrease while validation MSE (dashed) plateaus or diverges, confirming overfitting. The vertical dotted line marks the epoch of minimum validation loss.

hybrid comparison would require retraining the VQC without B, which was not performed. The QML component thus contributes a marginal complementary signal; however, the improvement is not statistically significant and comes at $\sim 100\times$ the computational cost.

3.4 AOP-Informed Features

Adding 6 binary MoA pathway features to the 17 molecular descriptors yields no improvement in RF LOO-CV performance (table 2).

Table 2: Effect of AOP-derived MoA features on RF LOO-CV performance.

Feature Set	N_{features}	**DegT50 R^2**	**K_{oc} R^2**
Molecular descriptors only	17	0.194	0.766
+ AOP MoA pathways	23	0.190	0.765
ΔR^2	+6	−0.004	−0.001

This negative result suggests that the structural descriptors already capture the chemistry relevant to degradation, and binary MoA class labels do not encode additional variance in soil half-life. We note that the remaining 40% of substances receive all-zero AOP vectors, adding noise without signal; however, restricting the analysis to the 67 covered substances yields comparable R^2 values, confirming that the null result is not a statistical artifact of sparse coverage.

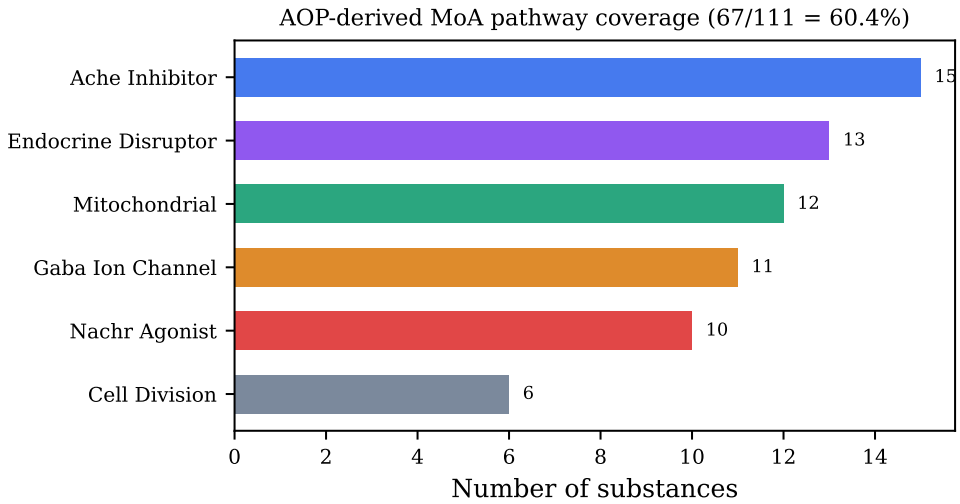

Figure 7: AOP-Wiki-derived mechanism-of-action pathway coverage across the 111 pesticide dataset. 67 substances (60.4%) map to at least one of 6 pathway categories.

4 Discussion

4.1 Why Does QML Fail on DegT50?

The negative R^2 of the VQC for DegT50 prediction is not a failure of quantum computing per se, but rather a predictable consequence of applying a 301-parameter model to 111 data points. Several factors contribute:

1. **Overfitting:** The $N_{\text{params}}/N_{\text{data}}$ ratio of 2.7 virtually guarantees memorization of training data. Classical models with far fewer effective parameters (RF uses bagging and feature subsampling as implicit regularization) are more robust.

2. **Barren plateaus:** With 12 qubits and 8 layers, the VQC may suffer from vanishing gradients in portions of the parameter landscape (McClean et al., 2018), making optimization unreliable with limited training data for validation.

3. **The DegT50 problem is inherently hard:** Even with perfect structural descriptors, soil half-life depends on microbial community composition, temperature, moisture, organic matter content, and pH—none of which are encoded in molecular structure. An R^2 ceiling near 0.2–0.3 may be fundamental for structure-only models.

4. **Linear readout bottleneck:** The VQC readout layer is a linear combination of $\langle Z_i \rangle$ expectation values. Any function the circuit can express must therefore lie in the span of these N_{qubits} observables. Given that Ridge regression—itself a linear model—already achieves $R^2 = 0.842$ for K_{oc} with the 17-feature set (inflated by bioaccessibility circularity; the fair 16-feature R^2 is 0.546), it is unclear whether the quantum feature map provides additional expressive power beyond what a linear model over the raw descriptors captures.

To disentangle these hypotheses, we conducted a systematic architecture search testing four circuit variants via 5-fold CV on DegT50 only (table 3): (A) a shallow-wide circuit (12 qubits, 3 layers, all-to-all CZ entanglement, double data re-uploading); (B) a deep-narrow circuit (8 qubits, 12 layers, ring CNOT, triple re-uploading); (C) a hardware-efficient template using PennyLane's `StronglyEntanglingLayers`; and (D) the baseline architecture with a reduced learning rate (0.03) and extended training (120 epochs).

Table 3: Architecture experiment: DegT50 prediction (5-fold CV) across four circuit variants. The current baseline achieves $R^2 = -0.141$ with 8 layers and 60 epochs.

Variant	Architecture	N_q	L	Epochs	Train MSE	R^2	Time
A	All-to-all CZ, 2× re-upload	12	3	80	0.183	−0.157	2.1 h
B	Ring CNOT, 3× re-upload	8	12	100	0.003	−0.190	3.3 h
C	`StronglyEntanglingLayers`	12	5	60	0.165	−0.146	1.3 h
D	Baseline+, lr=0.03	12	8	120	0.009	−0.205	4.1 h
Current baseline (lr=0.04)		12	8	60	—	−0.141	—

No variant improves upon the baseline. Remarkably, the two variants that achieve near-zero training MSE—variant B (0.003) and variant D (0.009)—produce the *worst* test R^2 (-0.190 and -0.205, respectively), while the underfit variants A and C (training MSE $\sim$0.17) perform marginally better. This anti-correlation between training fitness and generalization is a textbook signature of severe overfitting and demonstrates that the VQC has *sufficient* expressivity to memorize DegT50 training data but cannot extract a generalizable structure–activity relationship. The failure is therefore intrinsic to the target property and feature set, not to circuit architecture. Notably, variant A uses only 3 layers ($N_{\text{params}}/N_{\text{data}} \approx 1.0$), yet still produces negative R^2, demonstrating that reducing the parameter count alone does not rescue DegT50 prediction.

4.2 When Might QML Add Value?

Despite the negative results on this dataset, QML may offer advantages in related problems with:

- Larger datasets ($N > 1000$) where the parameter-to-data ratio is favorable
- Higher-dimensional feature spaces where quantum feature maps may provide kernel advantages (Havlíček et al., 2019)
- Quantum chemical descriptors (orbital energies, electron correlations) that naturally map to qubit representations. For example, the N_D electron correlation descriptor derived from coupled-cluster natural orbital occupations (Arda, 2026) quantifies multireference character across molecular systems and could serve as a qubit-native feature encoding, replacing or augmenting classical RDKit descriptors with information rooted in the electronic wavefunction.

4.3 The Value of Negative Results

This study contributes to the growing body of critical QML benchmarks that distinguish genuine quantum advantage from overstatement. Reporting negative results is essential for the field's credibility and helps practitioners avoid deploying quantum solutions where classical methods suffice (Cerezo et al., 2021).

5 Conclusions

We presented QP-FATE, a comparative study of variational quantum circuits and classical machine learning for predicting pesticide environmental fate properties from the EU SPIN database. Key findings:

1. **Gradient Boosting leads $\mathbf{K_{oc}}$ prediction:** On the circularity-free 16-feature set (excluding bioaccessibility B), GB achieves $R^2 = 0.764$ (LOO-CV), followed by RF (0.726). Linear models (Ridge 0.546, LASSO 0.535) are substantially weaker, indicating that K_{oc} benefits from nonlinear feature interactions. Notably, Ridge's previously reported $R^2 = 0.842$ was inflated by the circular inclusion of $B = \log_{10}(S/K_{oc})$, which contains the target variable.

2. **DegT50 remains unpredictable from structure:** All models achieve $R^2 \leq 0.22$ (LOO-CV), and a systematic architecture search across four VQC variants (table 3) confirms the failure is intrinsic to the target property, not circuit design.

3. **VQC underperforms classical baselines** on both properties, attributed to the unfavorable parameter-to-data ratio ($301/111 \approx 2.7$) and absence of regularization. Reducing the circuit to 3 layers ($N_{\text{params}}/N \approx 1$) does not recover predictive power.

4. **Hybrid stacking yields marginal gains;** the optimized blending weight $\alpha = 0.20$ improves RF R^2 by $\sim$0.02, but the improvement is not statistically significant and comes at $\sim$100$\times$ computational cost. AOP-derived MoA features provide no improvement, indicating that mechanism-of-action labels do not encode additional variance beyond structural descriptors.

5. **Minimum dataset recommendation:** Based on the observed failure at $N = 111$ and the N_{params}/N analysis, we estimate that VQC generalization for circuits of this size requires $N \gtrsim 1000$ samples—an order of magnitude beyond current regulatory databases.

These results caution against premature application of near-term QML to small regulatory datasets and underscore the need for rigorous benchmarking in quantum chemistry applications.

Acknowledgments

All quantum circuit simulations were performed using the PennyLane statevector simulator on consumer-grade hardware. Classical baseline computations used scikit-learn. The author thanks the PennyLane and scikit-learn developer communities for their open-source contributions.

Data Availability

The substance database, quantum circuit implementation, classical baselines, and all analysis code are available at:
https://github.com/unearthlyimprint/Quantum-Pesticide-Fate-Modeling

Conflicts of Interest

The author declares no competing interests.

References

Regulation (EC) No 1107/2009 of the European Parliament and of the Council concerning the placing of plant protection products on the market. *Official Journal of the European Union* **2009**, *L 309*, 1–50, OJ L 309, 24.11.2009.

FOCUS Groundwater Scenarios in the EU Review of Active Substances; 2006; Report of the FOCUS Work Group on Groundwater Scenarios, Version 2.

OECD Guideline for the Testing of Chemicals, Test No. 307: Aerobic and Anaerobic Transformation in Soil; 2002.

OECD Guideline for the Testing of Chemicals, Test No. 106: Adsorption – Desorption Using a Batch Equilibrium Method; 2000.

Breiman, L. Random Forests. *Machine Learning* **2001**, *45*, 5–32.

Friedman, J. H. Greedy function approximation: A gradient boosting machine. *Annals of Statistics* **2001**, *29*, 1189–1232.

Mansouri, K.; Grulke, C. M.; Judson, R. S.; Williams, A. J. OPERA models for predicting physicochemical properties and environmental fate endpoints. *Journal of Cheminformatics* **2018**, *10*.

Cerezo, M.; Arrasmith, A.; Babbush, R.; Benjamin, S. C.; Endo, S.; Fujii, K.; McClean, J. R.; Mitarai, K.; Yuan, X.; Cincio, L.; Coles, P. J. Variational quantum algorithms. *Nature Reviews Physics* **2021**, *3*, 625–644.

Havlíček, V.; Córcoles, A. D.; Temme, K.; Harrow, A. W.; Kandala, A.; Chow, J. M.; Gambetta, J. M. Supervised learning with quantum-enhanced feature spaces. *Nature* **2019**, *567*, 209–212.

Schuld, M.; Bocharov, A.; Svore, K. M.; Wiebe, N. Circuit-centric quantum classifiers. *Physical Review A* **2020**, *101*, 032308.

Pérez-Salinas, A.; Cervera-Lierta, A.; Gil-Fuster, E.; Latorre, J. I. Data re-uploading for a universal quantum classifier. *Quantum* **2020**, *4*, 226.

Schuld, M.; Sweke, R.; Meyer, J. J. Effect of data encoding on the expressive power of variational quantum machine-learning models. *Physical Review A* **2021**, *103*, 032430.

Bergholm, V.; Izaac, J.; Schuld, M.; Gogolin, C.; Ahmed, S.; Ajber, V.; Alam, M. S.; Alonso-Linaje, G.; AkashNarayanan, B.; Asber, A.; others PennyLane: Automatic differentiation of hybrid quantum-classical computations. *Quantum* **2022**, *6*, 660.

EU Substances, Preparations and Indicators Navigator (EU SPIN). https://food.ec.europa.eu/plants/pesticides/eu-pesticides-database_en, 2024; Accessed: 2025.

Kingma, D. P.; Ba, J. Adam: A Method for Stochastic Optimization. *Proceedings of the 3rd International Conference on Learning Representations (ICLR)* **2015**, arXiv:1412.6980.

Pedregosa, F.; Varoquaux, G.; Gramfort, A.; Michel, V.; Thirion, B.; Grisel, O.; Blondel, M.; Prettenhofer, P.; Weiss, R.; Dubourg, V.; others Scikit-learn: Machine Learning in Python. *Journal of Machine Learning Research* **2011**, *12*, 2825–2830.

Hoerl, A. E.; Kennard, R. W. Ridge Regression: Biased Estimation for Nonorthogonal Problems. *Technometrics* **1970**, *12*, 55–67.

Tibshirani, R. Regression Shrinkage and Selection via the Lasso. *Journal of the Royal Statistical Society: Series B* **1996**, *58*, 267–288.

AOP-Wiki: Adverse Outcome Pathway Knowledge Base. https://aopwiki.org/, 2024; Accessed: 2025.

Villeneuve, D. L.; Crump, D.; Garcia-Reyero, N.; Hecker, M.; Hutchinson, T. H.; LaLone, C. A.; Landesmann, B.; Lettieri, T.; Munn, S.; Nepelska, M.; others Adverse outcome pathway (AOP) development I: Strategies and principles. *Toxicological Sciences* **2014**, *142*, 312–320.

McClean, J. R.; Boixo, S.; Smelyanskiy, V. N.; Babbush, R.; Neven, H. Barren plateaus in quantum neural network training landscapes. *Nature Communications* **2018**, *9*, 4812.

Arda, C. A Unified Quantum Descriptor for Electron Correlation from Coupled-Cluster Natural Orbital Occupations. *Manuscript in preparation* **2026**, Manuscript in preparation.